


Agentic AI Workflow Ledger

I use AI agents the way a security team should use any privileged automation: least privilege, separation of duties, a human approval gate before anything changes, and an audit trail that records the agent's mistakes as carefully as its output.

This ledger covers twenty days of real work on my own projects: my public professional profiles, two open portfolio repositories, a job search, coursework and a home security lab. It is written for hiring managers and practitioners who want to see how I think about control, not just what the tools can do.

AT A GLANCE

The control system, by the numbers



METHOD

Five stages, with reading kept apart from changing

Each stage is a separate prompt template with its own permissions. An audit can't edit anything, a draft can't publish, and an apply step can only carry out what was approved.




The loop closes through the templates. When a run exposes a gap, the fix becomes a rule in the next template version, so later runs don't depend on anyone remembering the lesson. The main audit template reached version 20 in this period.

ACTIVITY

Twenty days of work, by day

Dated items created each day (Pacific time): workflow design revisions, commits to my public repositories, and the documents each workstream produced.



SECURITY POSTURE

Known agent risks, and the control for each

Mapped to the OWASP Top 10 for LLM Applications (2025) and the four functions of the NIST AI Risk Management Framework, so a reviewer can check the reasoning against recognized standards.

RISK	CONTROL IN THIS WORKFLOW	EVIDENCE KEPT	NIST AI RMF
Excessive agency OWASP LLM06:2025	Read, draft and apply are separate templates. Each public change needs its own approval. Job applications are reviewed and submitted by hand. The agent evaluates posts, comments, and messages and recommends next steps. Nothing is posted, sent, or changed without my explicit approval for that specific item, and the agent re-checks the live result afterward.	Dated "Applied" log for every change run	Manage
Prompt injection OWASP LLM01:2025	Audits are read-only, and nothing changes without per-item approval, so text planted in a web page can't trigger an action on its own.	Approval recorded before each apply step	Manage
Misinformation OWASP LLM09:2025	Every factual claim carries its source and the time it was read. Unverified items are labeled by strength rather than smoothed over. Skills backed only by exam objectives get knowledge-level wording, never "implemented."	Truth-check tables; error log below	Measure
Sensitive information disclosure OWASP LLM02:2025	No credentials, internal addresses or personal identifiers in any output. The agent never enters credentials. Where work involves regulated data such as patient records, that data never enters an AI tool.	Written working rules; redacted evidence only	Govern
Integrity and accountability Change management	Every file is versioned and time-stamped. Since Sep 26, superseded versions are archived rather than deleted, replacing an earlier trash-and-recreate practice. Each storage folder has one owning workflow, so parallel sessions can't overwrite each other.	Version history and archive folders	Govern
Stale or cached data Data quality	Audits must re-read live sources each run. A section that looks empty is re-checked with a real scroll before being reported, because pages that load on demand can appear blank.	Template rule added after a false report	Map

QUALITY CONTROL, WITH CONTEXT

Errors caught, and what changed because of them

I publish these because an AI workflow is only as trustworthy as its error handling. An agent that never seems to be wrong is one nobody is checking. Each entry led to a permanent change in the process.



Caught in internal review, before anything was published

Three errors of other kinds. The 12 false alarms have their own section and full log below.

DATE	WHAT HAPPENED	LASTING CHANGE	REACHED PUBLIC?
2026-09-16	An audit run suppressed verification steps it was supposed to perform	Verification scope written explicitly into the template	No
2026-09-23	An audit draft contained a fabricated citation	Sourcing failures escalated to elevated severity	No
2026-09-30	A browser editor mangled a multi-line paste	Caught before commit; full-file replacements done by upload	No

Found in my own public repositories by the audits, then fixed




DATE FOUND	WHAT THE AUDIT FOUND	FIX	STATUS
2026-09-18	A public README still held an internal note to confirm its sample data contained no personal information	Data confirmed non-personal; note removed	Fixed same day
2026-09-18	A public README pointed to a configuration file that didn't exist. Three audits flagged it, but it stayed open for five days.	A check against a fresh clone settled that the file was needed; it was created on Sep 23, and every referenced file is now checked for existence	Fixed Sep 23
2026-09-26	A vulnerability table in a public write-up had severity scores and CVE mappings that didn't match the National Vulnerability Database	Each entry re-checked against NVD and corrected; a likely false positive labeled as such	Fixed

The README item taught me the most. Three audits reported it, but each judged it from the rendered page, so no one could say for certain whether the file was needed, and it stayed open. Flagging a problem isn't the same as closing it. A mechanical check against a fresh copy of the repository settled it, and that check now runs every time.

VERIFICATION

False alarms: when a check says "no," confirm it before acting

An audit that reports something as missing, empty, broken or mismatched when it is fine is raising a false positive. Over twenty days the audits raised **12** of them. Every one was caught before it led to an unneeded change on a live profile, as far as the records show. Most were caught when I challenged the finding.



Every false positive on record



DATE	WHAT THE CHECK REPORTED	WHAT WAS TRUE	WHY IT WAS WRONG	CAUGHT BY
2026-09-30	Two job-site profiles essentially unbuilt	One already strong; the other held only duplicate stubs	First-pass judgment	Follow-up review
2026-09-30	Certificates on profile don't match my tracker	They matched	Hidden or partly loaded list	My challenge
2026-09-30	Education, Skills and Featured empty	All populated	Hidden or partly loaded list	My challenge
2026-09-28	Degree name differs between profile and resume	Wording matched	Misread or second-hand	My challenge
2026-09-28	Only 10 certificates listed; three missing	14 listed, including all three	Hidden or partly loaded list	The agent's own re-check
2026-09-26	A link fix not live	Live; the anonymous view was cached	Stale or cached copy	Not recorded
2026-09-23	A one-year date mismatch between profile and resume	No real mismatch	Hidden or partly loaded list	Not recorded
2026-09-23	A teaching stint missing	Present; detail was attributed to the wrong entry	Hidden or partly loaded list	Not recorded
2026-09-23	Three early-career roles missing	All three present in full	Hidden or partly loaded list	My challenge
2026-09-18	Most recent role shown as current job	The role had an end date	Misread or second-hand	My challenge
2026-09-18	Website typo fixes not live	Fixed; the audit read a cached copy	Stale or cached copy	My challenge
2026-09-18	Education section empty	It was populated; I checked my live profile directly	Hidden or partly loaded list	My challenge

The rule that came out of it: a finding that something is missing, empty, expired or mismatched stays unconfirmed until a second source agrees, ideally the primary one, read live. The lookup returned a false negative ("not there"), which raised a false-positive alarm. Confirming negatives before acting is what kept all of them harmless.

COMMUNICATION

How I write to the agent: courteous on hand-offs, specific on pushback

I brief the agent the way I'd brief a colleague. I counted my "please" and "thank you" across 4 project conversations and sorted each by the kind of message it came with. Courtesy is not a control; the controls are in the sections above. The count shows a working habit: polite when handing work off or accepting it, and still specific when pushing back.



PROJECT CONVERSATION	"PLEASE"	INSTRUCTIONS	APPROVALS	CHALLENGES THAT CAUGHT AN ERROR	CORRECTIONS	"THANK YOU"
Social media profiles	18	12	3	3	0	2
This portfolio project	11	6	2	0	3	2
Job search	13	13	0	0	0	8
Home security lab	2	2	0	0	0	0
All	44	33	5	3	3	12

How this was counted: each project's agent counted my messages in its own conversation, as of Oct 1, 2026. Most counts were rebuilt from condensed conversation records, so treat them as close, not exact: A message sent twice is counted once. The windows differ by project and some reach past the Sep 11–30 period used elsewhere on this page. The pattern matters more than the totals: not every challenge carried a "please," but every one was specific enough to check.

SCOPE

Where the agent helps, and where I keep it out

The agent does	The agent doesn't
- Read-only audits of public pages and repositories - Research, with every source logged and dated - Drafts: profile copy, documentation, cover letters and plans - Evaluate posts, comments, and messages and recommend next steps - Send approved comments, invitations and messages, one item at a time - Approved, individually initiated commits to my own repositories - Record-keeping: versioning, archiving and run logs	- Act without explicit approval for that specific item - Submit job applications; I review and submit each by hand - Enter passwords or handle credentials - Touch protected health information or client systems - Delete records; since Sep 26, superseded versions are archived

EVIDENCE

Public work you can inspect

home-cybersecurity-lab Isolated virtual lab with attacker and target machines. Includes a security policy, ethical-use terms and a license, all added through approved, logged runs.	quickstart-cybersecurity-career-sims Three bootcamp career simulations written up as capture-the-flag style reports, with the vulnerability table corrected against NVD.	github.com/ChowHusky21 Profile and commit history. 34 contributions in the last year as of Oct 6, 2026, up from 8 at the first audit on Sep 13.
--	--	---